

Project 2: Understanding Cache Memories

519021910366, Dou Yiming, douyiming@sjtu.edu.cn

June 5, 2021

1 Introduction

This project aims at making us understand cache memories better. The whole project can be divided into two parts.

Firstly, we are required to simulate a cache including sets, ways, tags and valid bits. It also covers cache replacement algorithm LRU that we have learned in operating system class. I complete this part with reference to slides of the course and CSAPP book.

Secondly, on the basis of the cache simulator, we are required to implement an optimization. When transposing the matrix, the cache may not play to its full capability if without enough design. Therefore, we need to improve the cache hit rate by blocking the matrix and transfer each block with the help of cache in order to speed up.

2 Experiments

2.1 Part A

2.1.1 Analysis

The cache we are required to simulate contains 2^s sets, E ways and 2^b block offsets, whose structure can be shown in Fig. 1.

In the process of a single entry, we begin with traversing each set, checking whether it's valid or not and whether the tag is the same as the data we want. If a qualified set is found, we define this as a cache hit. Otherwise, if the cache is full, we define this situation as eviction, in which we should swap a new element with the existing element in the cache.

The next problem is how we choose the place we want to evict. In this project, we use the Least-Recently-Used (LRU) strategy, meaning that we need to seek the data that we have not used for the longest time. Therefore, we attach a time stamp for each set, each time the set is visited, we update the time stamp of it to the current time, and when eviction happens, we choose the set whose time stamp is the smallest.

Four-Way Set Associative Cache

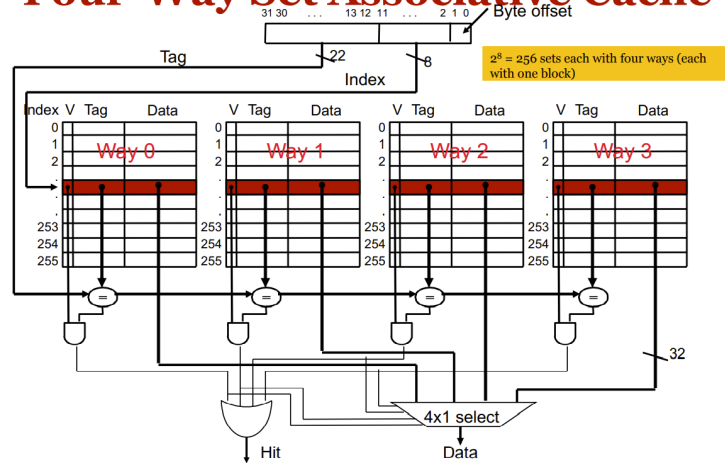


Figure 1: Cache Structure

2.1.2 Code

```

1  /*
2   * Name: DouYiming
3   * ID: 519021910366
4   */
5  #include <getopt.h>
6  #include <stdlib.h>
7  #include <stdio.h>
8  #include "cachelab.h"
9
10 const int m = 64;
11
12 typedef struct Arg
13 {
14     int s; // Number of set index bits
15     int E; // number of lines per set
16     int b; // Number of block bits
17     FILE *f; // valgrind trace to replay
18 } Arg;
19
20 typedef struct summary
21 {
22     int hits;
23     int misses;
24     int evictions;
25 } summary;
26
27 typedef struct Line
28 {
29     int valid;
30     int tag;
31     int time_stamp;

```

```

32 } Line;
33
34 typedef struct Set
35 {
36     Line* lines;
37 } Set;
38
39 Arg arg_parser(int argc, char ** argv)
40 {
41     Arg arg;
42     if(getopt(argc, argv, "s:E:b:t") == 's')
43         arg.s = atoi(optarg);
44     if(getopt(argc, argv, "s:E:b:t") == 'E')
45         arg.E = atoi(optarg);
46     if(getopt(argc, argv, "s:E:b:t") == 'b')
47         arg.b = atoi(optarg);
48     if(getopt(argc, argv, "s:E:b:t") == 't')
49         arg.f = fopen(optarg, "r");
50     return arg;
51 }
52
53 Set* init_cache(Arg arg)
54 {
55     int set_num = 1 << arg.s;
56     Set* caches=(Set*)malloc(set_num * sizeof(Set));
57     for(int i = 0; i < set_num; ++i)
58     {
59         caches[i].lines = (Line*)malloc(arg.E * sizeof(Line));
60         for(int j = 0; j < arg.E; ++j)
61         {
62             caches[i].lines[j].valid = 0;
63             caches[i].lines[j].time_stamp = 0;
64         }
65     }
66     return caches;
67 }
68
69 summary_simulator(Arg arg, Set* caches)
70 {
71     // information to display
72     summary ans;
73     ans.hits = 0;
74     ans.misses = 0;
75     ans.evictions = 0;
76     // compute tag bits
77     int t = m - arg.b - arg.s;
78     // simulator time
79     int cur_time = 0;
80     // information of each instruction
81     char op;
82     long addr;
83     int len;
84     int cycle_time;
85
86     // read from file
87     while(!feof(arg.f))
88     {

```

```

89     if (fscanf(arg.f, " %c %lx,%d", &op, &addr, &len) != 3)
90         continue;
91
92     if (op == 'L' || op == 'S')
93         cycle_time = 1;
94     else if (op == 'M')
95         cycle_time = 2;
96     else
97         continue;
98     while(cycle_time--)
99     {
100         ++cur_time;
101         int set_num = (addr >> arg.b) & ((1 << arg.s) - 1);
102         int tag = (addr >> (arg.s + arg.b)) & ((1 << t) - 1);
103         int is_hit = 0;
104         int invalid_line = -1;
105         int LRU_line = -1;
106         int min_time = 0x7fffffff;
107         // search each way
108         for(int i = 0; i < arg.E; ++i)
109         {
110             Line* line = & caches[set_num].lines[i];
111             if(line->valid && line->tag == tag) // hit
112             {
113                 is_hit = 1;
114                 line->time_stamp = cur_time;
115                 break;
116             }
117             else if(!line->valid)
118             {
119                 invalid_line = i;
120                 break;
121             }
122             else if(line->valid && line->time_stamp < min_time)
123             {
124                 min_time = line->time_stamp;
125                 LRU_line = i;
126             }
127         }
128         if(is_hit) // hit
129             ++ans.hits;
130         else
131         {
132             ++ans.misses;
133             if(invalid_line != -1) // miss
134             {
135                 Line* l = & caches[set_num].lines[invalid_line];
136                 l->tag = tag;
137                 l->time_stamp = cur_time;
138                 l->valid = 1;
139             }
140             else // evict
141             {
142                 ++ans.evictions;
143                 Line* l = & caches[set_num].lines[LRU_line];
144                 l->tag = tag;
145                 l->time_stamp = cur_time;

```

```

146         }
147     }
148 }
149 }
150 return ans;
151 }
152
153 void terminate(Arg arg, Set* caches)
154 {
155     fclose(arg.f);
156     for (int i = 0; i < (1 << arg.s); ++i)
157         free(caches[i].lines);
158     free(caches);
159 }
160
161 int main(int argc, char **argv)
162 {
163     // arg parsing
164     Arg arg = arg_parser(argc, argv);
165     // cache initialization
166     Set* caches = init_cache(arg);
167     // simulation
168     summary s = simulator(arg, caches);
169     // ending
170     terminate(arg, caches);
171     // print ans
172     printSummary(s.hits, s.misses, s.evictions);
173     return 0;
174 }

```

```

dodo@ubuntu:~/code/Computer_Architecture/project2-handout$ ./test-csim
Points (s,E,b)    Hits    Misses    Evicts    Hits    Misses    Evicts
3 (1,1,1)         9         8         6         9         8         6 traces/yi2.trace
3 (4,2,4)         4         5         2         4         5         2 traces/yi.trace
3 (2,1,4)         2         3         1         2         3         1 traces/dave.trace
3 (2,1,3)        167        71        67        167        71        67 traces/trans.trace
3 (2,2,3)       201        37        29        201        37        29 traces/trans.trace
3 (2,4,3)       212        26        10        212        26        10 traces/trans.trace
3 (5,1,5)       231         7         0        231         7         0 traces/trans.trace
6 (5,1,5)    265189    21775    21743    265189    21775    21743 traces/long.trace
27
TEST_CSIM_RESULTS=27

```

Figure 2: The result of part A

2.1.3 Evaluation

The result of part A is shown in Fig. 2.

2.2 Part B

2.2.1 Analysis

In this part, we need to optimize the process of transposing matrix of different sizes.

The basic idea is that we should use the elements in the cache as much as possible for each time. Therefore, simply traversing the rows and columns of a large matrix may not be a good idea due to large amounts of evictions. In order to make full use of the cache, we separate the matrix into blocks and move a block into the cache each time, ensuring that most elements in cache are used.

Furthermore, our method to separate the matrix should vary when facing different sizes of matrix by calculating the address of elements in the matrix. Therefore, we respectively implement a optimization method for each matrix transposing situation.

2.2.2 Code

```

1  /*
2  * Name: Dou Yi-ming
3  * ID: 519021910366
4  * trans.c — Matrix transpose B = A^T
5  *
6  * Each transpose function must have a prototype of the form:
7  * void trans(int M, int N, int A[N][M], int B[M][N]);
8  *
9  * A transpose function is evaluated by counting the number of
10 * misses
11 */
12 #include <stdio.h>
13 #include "cachelab.h"

```

```

14
15 int is_transpose(int M, int N, int A[N][M], int B[M][N]);
16
17 /*
18  * transpose_submit — This is the solution transpose function that
19  *   you
20  *   will be graded on for Part B of the assignment. Do not
21  *   change
22  *   the description string "Transpose submission", as the driver
23  *   searches for that string to identify the transpose function
24  *   to
25  *   be graded.
26  */
27 char transpose_submit_desc[] = "Transpose submission";
28 void transpose_submit(int M, int N, int A[N][M], int B[M][N])
29 {
30     // local variables
31     int ii, jj, i, j;
32     int tmp0, tmp1, tmp2, tmp3, tmp4, tmp5, tmp6, tmp7;
33     if (M == 32 && N == 32)
34     {
35         // 8 * 8 blocks
36         for(ii = 0; ii < 4; ++ii)
37             for(jj = 0; jj < 4; ++jj)
38                 for(i = 0; i < 8; ++i)
39                 {
40                     for(j = 0; j < 8; ++j)
41                         if(i != j)
42                             B[8 * jj + j][8 * ii + i] = A[8 * ii +
43                             i][8 * jj + j];
44                     B[8 * jj + i][8 * ii + i] = A[8 * ii + i][8 *
45                     jj + i];
46                 }
47     }
48     else if (M == 64 && N == 64)
49     {
50         for(ii = 0; ii < 8; ++ii)
51             for(jj = 0; jj < 8; ++jj)
52             {
53                 for(i = 0; i < 4; ++i)
54                 {
55                     tmp0 = A[ii * 8 + i][jj * 8 + 0];
56                     tmp1 = A[ii * 8 + i][jj * 8 + 1];
57                     tmp2 = A[ii * 8 + i][jj * 8 + 2];
58                     tmp3 = A[ii * 8 + i][jj * 8 + 3];
59                     tmp4 = A[ii * 8 + i][jj * 8 + 4];
60                     tmp5 = A[ii * 8 + i][jj * 8 + 5];
61                     tmp6 = A[ii * 8 + i][jj * 8 + 6];
62                     tmp7 = A[ii * 8 + i][jj * 8 + 7];
63
64                     B[jj * 8 + 0][ii * 8 + i] = tmp0;
65                     B[jj * 8 + 1][ii * 8 + i] = tmp1;
66                     B[jj * 8 + 2][ii * 8 + i] = tmp2;
67                     B[jj * 8 + 3][ii * 8 + i] = tmp3;
68
69                     B[jj * 8 + 0][ii * 8 + i + 4] = tmp4;
70                     B[jj * 8 + 1][ii * 8 + i + 4] = tmp5;

```

```

66         B[jj * 8 + 2][ii * 8 + i + 4] = tmp6;
67         B[jj * 8 + 3][ii * 8 + i + 4] = tmp7;
68     }
69     for(j = 0; j < 4; ++j)
70     {
71         tmp0 = B[jj * 8 + j][ii * 8 + 4];
72         tmp1 = B[jj * 8 + j][ii * 8 + 5];
73         tmp2 = B[jj * 8 + j][ii * 8 + 6];
74         tmp3 = B[jj * 8 + j][ii * 8 + 7];
75
76         tmp4 = A[ii * 8 + 4][jj * 8 + j];
77         tmp5 = A[ii * 8 + 5][jj * 8 + j];
78         tmp6 = A[ii * 8 + 6][jj * 8 + j];
79         tmp7 = A[ii * 8 + 7][jj * 8 + j];
80
81         B[jj * 8 + j][ii * 8 + 4] = tmp4;
82         B[jj * 8 + j][ii * 8 + 5] = tmp5;
83         B[jj * 8 + j][ii * 8 + 6] = tmp6;
84         B[jj * 8 + j][ii * 8 + 7] = tmp7;
85
86         tmp4 = A[ii * 8 + 4][jj * 8 + j + 4];
87         tmp5 = A[ii * 8 + 5][jj * 8 + j + 4];
88         tmp6 = A[ii * 8 + 6][jj * 8 + j + 4];
89         tmp7 = A[ii * 8 + 7][jj * 8 + j + 4];
90
91         B[jj * 8 + j + 4][ii * 8 + 0] = tmp0;
92         B[jj * 8 + j + 4][ii * 8 + 1] = tmp1;
93         B[jj * 8 + j + 4][ii * 8 + 2] = tmp2;
94         B[jj * 8 + j + 4][ii * 8 + 3] = tmp3;
95         B[jj * 8 + j + 4][ii * 8 + 4] = tmp4;
96         B[jj * 8 + j + 4][ii * 8 + 5] = tmp5;
97         B[jj * 8 + j + 4][ii * 8 + 6] = tmp6;
98         B[jj * 8 + j + 4][ii * 8 + 7] = tmp7;
99     }
100 }
101 }
102 else if (M == 61 && N == 67)
103 {
104     for(ii = 0; ii < 5; ++ii)
105         for(jj = 0; jj < 4; ++jj)
106             for(i = 0; i < 16 && ii * 16 + i < N; ++i)
107                 for(j = 0; j < 16 && jj * 16 + j < M; ++j)
108                     B[16 * jj + j][16 * ii + i] = A[16 * ii + i
109 ] [16 * jj + j];
110 }
111 else
112 {
113     int tmp;
114     for (i = 0; i < N; i++)
115         for (j = 0; j < M; j++)
116         {
117             tmp = A[i][j];
118             B[j][i] = tmp;
119         }
120 }
121 }

```



```

122
123 /*
124  * You can define additional transpose functions below. We've
125  * defined
126  * a simple one below to help you get started.
127  */
128
129 /*
130  * trans — A simple baseline transpose function, not optimized for
131  * the cache.
132  */
133 char trans_desc[] = "Simple row-wise scan transpose";
134 void trans(int M, int N, int A[N][M], int B[M][N])
135 {
136     int i, j, tmp;
137     for (i = 0; i < N; i++) {
138         for (j = 0; j < M; j++) {
139             tmp = A[i][j];
140             B[j][i] = tmp;
141         }
142     }
143 }
144
145 /*
146  * registerFunctions — This function registers your transpose
147  * functions with the driver. At runtime, the driver will
148  * evaluate each of the registered functions and summarize
149  * their
150  * performance. This is a handy way to experiment with
151  * different
152  * transpose strategies.
153  */
154 void registerFunctions()
155 {
156     /* Register your solution function */
157     registerTransFunction(transpose_submit, transpose_submit_desc);
158
159     /* Register any additional transpose functions */
160     registerTransFunction(trans, trans_desc);
161 }
162
163 /*
164  * is_transpose — This helper function checks if B is the transpose
165  * of
166  * A. You can check the correctness of your transpose by
167  * calling
168  * it before returning from the transpose function.
169  */
170 int is_transpose(int M, int N, int A[N][M], int B[M][N])
171 {
172     int i, j;
173     for (i = 0; i < N; i++) {
174         for (j = 0; j < M; ++j) {

```

```
173         if (A[i][j] != B[j][i]) {  
174             return 0;  
175         }  
176     }  
177 }  
178 return 1;  
179 }
```

```
Part B: Testing transpose function
Running ./test-trans -M 32 -N 32
Running ./test-trans -M 64 -N 64
Running ./test-trans -M 61 -N 67
```

Cache Lab summary:			
	Points	Max pts	Misses
Csim correctness	27.0	27	
Trans perf 32x32	8.0	8	287
Trans perf 64x64	8.0	8	1187
Trans perf 61x67	10.0	10	1992
Total points	53.0	53	

Figure 3: Result of part B

2.2.3 Evaluation

The result of the test is shown in Fig. 3

3 Conclusion

3.1 Problems

1. The name of each part of the cache is different from the name we use in class and slides, which cause lots of confusion. With reference to CSAPP book, we finally solve this problem.
2. The optimization of the matrix transposing problem is far from our intuition, which is simply traversing the rows and columns. After referring to lots of materials, we eventually understand the separation method.

3.2 Achievements

1. My results are excellent. All of the requirements have been satisfied and i get full score in each test.
2. When designing the cache simulator, i paid a lot of attention to the structure of the code, and optimize the structure to improve readability and conciseness for a long time. Finally, i reduced the length of main function to 6 lines and i am very proud of it.
3. I figured out and successfully implemented the optimization method of transposing the matrix, which is very complicated.